

```
In [1]: import pandas as pd
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline
matplotlib inline
```

Generate 20 data pairs (X, Y) using  $y = \sin(2\pi x) + N$

Use uniform distribution between 0 and 1 for X

Sample N from the normal distribution

Use 10 for train and 10 for test

```
In [2]: p1 = np.pi
x = np.random.uniform(0,1,20)
n = np.random.normal(0,1)
```

```
In [3]: x.sort()
x
```

```
Out[3]: array([0.11556161, 0.1381175, 0.21395482, 0.25098457, 0.28618993,
0.32111972, 0.36554996, 0.38517852, 0.42391901, 0.55918,
0.63300403, 0.67981729, 0.71032217, 0.7622199, 0.77789733,
0.88015235, 0.80185753, 0.83039906, 0.90886866, 0.95325081])
```

```
In [4]: n
```

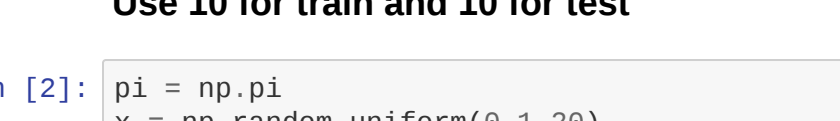
```
Out[4]: 0.8306261964867722
```

```
In [5]: y = (np.sin(2*pi*x) + n)
```

```
In [6]: y
```

```
Out[6]: array([ 1.48429892, 1.59372052, 1.80508937, 1.83061304, 1.80488474,
1.4820272, 1.0784764, 1.4911396, 1.2962617, 0.50393376,
0.8656869, -0.87397656, -0.126872, -0.1664767, 0.15395206,
-0.12013407, -0.11836602, -0.04446979, 0.28933629, 0.54121371])
```

```
In [7]: plt.scatter(x, y)
plt.plot(x, y, color = 'r')
plt.grid(True)
plt.show()
```



```
In [8]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.5)
```

Conversion of nd array to lists in Python

```
In [9]: X_train_list = X_train.tolist()
X_test_list = X_test.tolist()
y_train_list = y_train.tolist()
y_test_list = y_test.tolist()
```

Conversion of lists to dictionary

```
In [10]: dictionary1 = {}
for key in y_train_list:
    for value in X_train_list:
        dictionary1[key] = value
        y_train_list.remove(value)
        break
print(dictionary1)
```

```
{0.80185753177852: -0.118366021519314, 0.2139548237480998: 1.805089376896334, 0.1155616129890326: 1.484298921898674, 0.55918004892334: 0.50393376263384, 1.77789732928181, 0.830399063626938: -0.044469791762804, 0.633004032824374, 0.8801523518913915, 0.286189930079667: 1.804884739755416, 0.80015235177852: -0.120134071519315, 0.47901724627637: -0.0739765693932376}
```

```
In [11]: dictionary2 = {}
for key in X_train_list:
    for value in y_train_list:
        dictionary2[key] = value
        X_train_list.remove(value)
        break
print(dictionary2)
```

```
{0.4239091848272982: 1.296261745334545, 0.762219936017758: -0.1664767865946391, 0.13811757208734: 1.593720517999259, 0.77789733458884: -0.153952061280126, 0.38517853831859: 1.491139617929234, 0.305499778012032: 1.0784763872324, 0.953250819782385: 0.8412137113915, 0.8289045728678416: 1.8306261964867726, 0.55918044877489: 0.7503622747653951, -0.169371999415947, 0.968898078955147, 0.289336294883193}
```

Creating sorted lists

```
In [12]: sort_order1 = sorted(dictionary1.items())
sort_order1
```

```
Out[12]: [(0.11556161298908535, 1.484298921898674),
(0.2139548237480998, 1.805089376896334),
(0.2861899292976672, 1.804884739755416),
(0.321119726263364, 1.77789732928181),
(0.36554995729672, 1.0784763872324),
(0.38517853831859, 1.491139617929234),
(0.4239091848272982, 1.296261745334545),
(0.55918004892334, 0.50393376263384),
(0.633004032824374, 0.865686940684262),
(0.762219936017758, -0.1664767865946391),
(0.77789733458884, 0.153952061280126),
(0.78901724627637, -0.0739765693932376),
(0.80185753177852, -0.118366021519314),
(0.830399063626938, 0.153952061280126),
(0.8801523518913915, 0.289336294883193),
(0.908898078955147, 0.289336294883193),
(0.953250819782385, 0.5412137113915)]
```

```
In [13]: sort_order2 = sorted(dictionary2.items())
sort_order2
```

```
Out[13]: [(0.11556161298908535, 1.484298921898674),
(0.2139548237480998, 1.805089376896334),
(0.2509845728794916, 1.8306261964867726),
(0.2861899292976672, 1.804884739755416),
(0.38517853831859, 1.491139617929234),
(0.4239091848272982, 1.296261745334545),
(0.55918004892334, 0.50393376263384),
(0.762219936017758, -0.1664767865946391),
(0.77789733458884, 0.153952061280126),
(0.78901724627637, -0.0739765693932376),
(0.80185753177852, -0.118366021519314),
(0.830399063626938, 0.153952061280126),
(0.908898078955147, 0.289336294883193),
(0.953250819782385, 0.5412137113915)]
```

```
In [14]: sorted_x = []
sorted_y = []
for i in sort_order1:
    sorted_x.append(i[0])
    sorted_y.append(i[1])
```

```
In [15]: sorted_x2 = []
sorted_y2 = []
for i in sort_order2:
    sorted_x2.append(i[0])
    sorted_y2.append(i[1])
```

Conversion of lists to Arrays

```
In [16]: X_train_sorted = np.array(sorted_x)
Y_train_sorted = np.array(sorted_y)
X_test_sorted = np.array(sorted_x2)
Y_test_sorted = np.array(sorted_y2)
```

Created a Graph with X, Y and Y-Train

```
In [17]: plt.scatter(X_train_sorted, Y_train_sorted, color = 'b', label = 'Train Data')
plt.plot(x, y, color = 'g', label = 'Total Data')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [18]: plt.scatter(X_test_sorted, Y_test_sorted, color = 'b', label = 'Test Data')
plt.plot(x, y, color = 'g', label = 'Total Data')
plt.legend()
plt.grid(True)
plt.show()
```



Using root mean square error, find weights of polynomial regression for order is 0, 1, 3, 9

Creating weights

```
In [19]: degrees = [0,1,3,9]
for degree in (degrees):
    W[degree] = (np.polyfit(X_train_sorted, Y_train_sorted, degree)).tolist()
print(W)
```

```
{0: [0.7086702324932741, 1: [-3.123045392348336, 2: 3.2427189772447, 3: [29.1619592875916, 4: -42.260649999863, 5: 15.15123082082819, 6: 2.204519969030498], 9: [-37.8320260151183, 106.805873252495, -251.8909746032773, 122.3287898951316, 26.48495972504952, 18.08384889806652, 4.740385518992494, 0.4353243075647953, 6.251464949493874, 0.8316174081854233]}
```

Display weights in tabular format

```
In [20]: df = pd.DataFrame.from_dict(W, orient='index')
df = df.fillna(0)
df = df.rename(columns={0: 'W0', 1: 'W1', 3: 'W3', 9: 'W9'})
df = df.rename({'W0': 'w0', 'W1': 'w1', 'W3': 'w3', 'W9': 'w9', '0': 'w0', '1': 'w1', '3': 'w3', '9': 'w9'})
df
```

```
Out[20]:
```

	W=0	W=1	W=3	W=9
w0	0.70867	1.123045	29.161959	122.328789
w1	0.00000	2.34267	-42.260649	106.805873
w3	0.00000	0.00000	15.151231	-251.890975
w9	0.00000	0.00000	0.00000	12.083849
w0	0.00000	0.00000	0.00000	36.40096
w1	0.00000	0.00000	0.00000	-44.74026
w3	0.00000	0.00000	0.00000	0.430324
w9	0.00000	0.00000	0.00000	6.251465
w0	0.00000	0.00000	0.00000	0.831617

Draw a chart of fit data


For M=0

```
In [21]: lin_reg = LinearRegression()
```

```
In [22]: poly0=PolynomialFeatures(degree=0)
x_poly0 = poly0.fit_transform(X_train_sorted.reshape(-1,1))
poly0.fit(x_train_sorted.reshape(-1,1),Y_train_sorted)
lin_reg.fit(x_poly0,Y_train_sorted)
```

```
Out[22]: LinearRegression()
```

```
In [23]: plt.scatter(X_train,y_train,color='blue',facecolors='none',edgecolors='b', label = 'Train Data')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(X_train_sorted,lin_reg.predict(poly0.fit_transform(X_train_sorted.reshape(-1,1))),color='red', label = 'M=0')
plt.legend()
plt.grid(True)
plt.show()
```




For M=1

```
In [24]: poly1=PolynomialFeatures(degree=1)
x_poly1 = poly1.fit_transform(X_train_sorted.reshape(-1,1))
poly1.fit(x_train_sorted.reshape(-1,1),Y_train_sorted)
lin_reg.fit(x_poly1,Y_train_sorted)
```

```
Out[24]: LinearRegression()
```

```
In [25]: plt.scatter(X_train,y_train,color='blue',facecolors='none',edgecolors='b', label = 'Train Data')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(X_train_sorted,lin_reg.predict(poly1.fit_transform(X_train_sorted.reshape(-1,1))),color='red', label = 'M=1')
plt.legend()
plt.grid(True)
plt.show()
```




For M=3

```
In [26]: poly3=PolynomialFeatures(degree=3)
x_poly3 = poly3.fit_transform(X_train_sorted.reshape(-1,1))
poly3.fit(x_train_sorted.reshape(-1,1),Y_train_sorted)
lin_reg.fit(x_poly3,Y_train_sorted)
```

```
Out[26]: LinearRegression()
```

```
In [27]: plt.scatter(X_train,y_train,color='blue',facecolors='none',edgecolors='b', label = 'Train Data')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(X_train_sorted,lin_reg.predict(poly3.fit_transform(X_train_sorted.reshape(-1,1))),color='red', label = 'M=3')
plt.legend()
plt.grid(True)
plt.show()
```

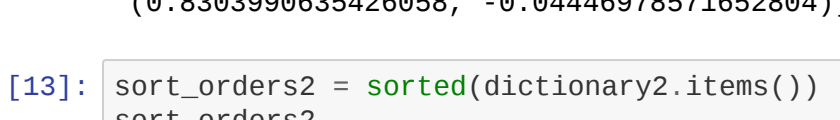


For M=9

```
In [28]: poly9=PolynomialFeatures(degree=9)
x_poly9 = poly9.fit_transform(X_train_sorted.reshape(-1,1))
poly9.fit(x_train_sorted.reshape(-1,1),Y_train_sorted)
lin_reg.fit(x_poly9,Y_train_sorted)
```

```
Out[28]: LinearRegression()
```

```
In [29]: plt.scatter(X_train,y_train,color='blue',facecolors='none',edgecolors='b', label = 'Train Data')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(X_train_sorted,lin_reg.predict(poly9.fit_transform(X_train_sorted.reshape(-1,1))),color='red', label = 'M=9')
plt.legend()
plt.grid(True)
plt.show()
```



Draw train error vs test error

```
In [30]: train_error = np.empty(10)
test_error = np.empty(10)
```

```
In [31]: for degree in range(10):
    est = make_pipeline(PolynomialFeatures(degree), LinearRegression())
    est.fit(X_train_sorted.reshape(-1,1),Y_train_sorted)
    train_error[degree] = np.sqrt(mean_squared_error(Y_train_sorted, est.predict(X_train_sorted.reshape(-1,1))))
    test_error[degree] = np.sqrt(mean_squared_error(Y_test_sorted, est.predict(X_test_sorted.reshape(-1,1))))
```

```
In [32]: train_error
```

```
Out[32]: array([0.81964873, 0.77259877, 0.82362493, 0.73728337, 0.80498832,
1.11634695, 1.12552523, 1.1261745, 1.21977385, 1.2234742,
1.29397878, 1.28907182, 1.23240875, 1.3444892, 1.38667388,
2.43317309, 1.42733959, 1.40383652, 1.40733982, 1.40880865,
1.58995293, 1.51226268, 1.52159494, 1.5474525, 1.5452586,
1.57688772, 1.58898901, 1.59688831, 1.63252593, 1.5551751,
1.52862343, 1.52473884, 1.51651545, 1.51271117, 1.4759718,
1.4748469, 1.37899235, 1.2751737, 1.2452644, 1.3408711,
1.3980921, 1.2785729, 1.38292786, 1.2364674, 1.2968337,
1.2550691, 1.1877663, 0.89497483, 0.8569232, 0.77989179,
0.76880007, 0.7211979, 0.7182026, 0.70897437, 0.6466539,
0.6425849, 0.5895814, 0.59537879, 0.39284785, 0.30775939,
0.15944467, 0.11737044, 0.11293362, 0.09062844, 0.0518779,
-0.25372644, -0.2575345, -0.28466944, -0.25167973, 0.37759939,
-0.369362, -0.3744227, -0.3855121, -0.39945219, -0.3923249,
-0.429711, -0.42898903, -0.43148438, -0.43186033, 0.48879729,
-0.40191223, -0.39899397, -0.41203583, -0.2540812, 0.2489532,
-0.15735692, -0.8153487, -0.87809585, -0.8489618, 0.48572097,
0.55585153, 0.87441225, 0.28703023, 0.24584929, 0.28846028,
0.2811844, 0.38892113, 0.34892965, 0.34646866, 0.47323571])
```

```
In [33]: n2
```

```
Out[33]: 0.568601979303957
```

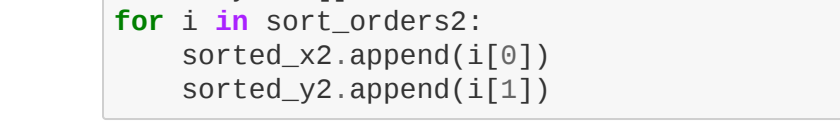
```
In [38]: y2 = (np.sin(2*pi*x2) + n2)
```

```
Out[38]: array([0.6914873, 0.77259877, 0.82362493, 0.73728337, 0.80498832,
1.11634695, 1.12552523, 1.1261745, 1.21977385, 1.2234742,
1.29397878, 1.28907182, 1.23240875, 1.3444892, 1.38667388,
2.43317309, 1.42733959, 1.40383652, 1.40733982, 1.40880865,
1.58995293, 1.51226268, 1.52159494, 1.5474525, 1.5452586,
1.57688772, 1.58898901, 1.59688831, 1.63252593, 1.5551751,
1.52862343, 1.52473884, 1.51651545, 1.51271117, 1.4759718,
1.4748469, 1.37899235, 1.2751737, 1.2452644, 1.3408711,
1.3980921, 1.2785729, 1.38292786, 1.2364674, 1.2968337,
1.2550691, 1.1877663, 0.89497483, 0.8569232, 0.77989179,
0.76880007, 0.7211979, 0.7182026, 0.70897437, 0.6466539,
0.6425849, 0.5895814, 0.59537879, 0.39284785, 0.30775939,
0.15944467, 0.11737044, 0.11293362, 0.09062844, 0.0518779,
-0.25372644, -0.2575345, -0.28466944, -0.25167973, 0.37759939,
-0.369362, -0.3744227, -0.3855121, -0.39945219, -0.3923249,
-0.429711, -0.42898903, -0.43148438, -0.43186033, 0.48879729,
-0.40191223, -0.39899397, -0.41203583, -0.2540812, 0.2489532,
-0.15735692, -0.8153487, -0.87809585, -0.8489618, 0.48572097,
0.55585153, 0.87441225, 0.28703023, 0.24584929, 0.28846028,
0.2811844, 0.38892113, 0.34892965, 0.34646866, 0.47323571])
```

```
In [39]: x2_poly9 = poly9.fit_transform(x2.reshape(-1,1))
poly9.fit(x2.reshape(-1,1),y2)
lin_reg.fit(x2_poly9,y2)
```

```
Out[39]: LinearRegression()
```

```
In [40]: plt.scatter(x2,y2,color='blue',facecolors='none',edgecolors='b', label = 'M=180')
plt.plot(x2,y2,color='green', label = 'Y=x+1, X')
plt.plot(x2,lin_reg.predict(poly9.fit_transform(x2.reshape(-1,1))),color='red', label = 'M=9')
plt.title('M=9')
plt.legend()
plt.grid(True)
plt.show()
```




Now we will regularize using the sum of weights.

Draw chart for lambda is 1, 1/10, 1/100, 1/1000, 1/10000

```
In [41]: poly_features = PolynomialFeatures(degree = 9)
ridge = Ridge(alpha = 1)
pip = Pipeline([('poly_features', poly_features),('ridge', ridge)])
pip.fit(x.reshape(-1,1),y)
a = pip.predict(x.reshape(-1,1))
```

```
plt.scatter(x,y,color='blue',facecolors='none',edgecolors='b')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(x,a,color='red', label = 'lambda')
plt.title('lambda = 1')
plt.legend()
plt.grid(True)
plt.show()
```




```
In [42]: poly_features = PolynomialFeatures(degree = 9)
ridge = Ridge(alpha = 1/10)
pip = Pipeline([('poly_features', poly_features),('ridge', ridge)])
pip.fit(x.reshape(-1,1),y)
a = pip.predict(x.reshape(-1,1))
```

```
plt.scatter(x,y,color='blue',facecolors='none',edgecolors='b')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(x,a,color='red', label = 'lambda')
plt.title('lambda = 1/10')
plt.legend()
plt.grid(True)
plt.show()
```




```
In [43]: poly_features = PolynomialFeatures(degree = 9)
ridge = Ridge(alpha = 1/100)
pip = Pipeline([('poly_features', poly_features),('ridge', ridge)])
pip.fit(x.reshape(-1,1),y)
a = pip.predict(x.reshape(-1,1))
```

```
plt.scatter(x,y,color='blue',facecolors='none',edgecolors='b')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(x,a,color='red', label = 'lambda')
plt.title('lambda = 1/100')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [44]: poly_features = PolynomialFeatures(degree = 9)
ridge = Ridge(alpha = 1/1000)
pip = Pipeline([('poly_features', poly_features),('ridge', ridge)])
pip.fit(x.reshape(-1,1),y)
a = pip.predict(x.reshape(-1,1))
```

```
plt.scatter(x,y,color='blue',facecolors='none',edgecolors='b')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(x,a,color='red', label = 'lambda')
plt.title('lambda = 1/1000')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [45]: poly_features = PolynomialFeatures(degree = 9)
ridge = Ridge(alpha = 1/10000)
pip = Pipeline([('poly_features', poly_features),('ridge', ridge)])
pip.fit(x.reshape(-1,1),y)
a = pip.predict(x.reshape(-1,1))
```

```
plt.scatter(x,y,color='blue',facecolors='none',edgecolors='b')
plt.plot(x,y,color='green', label = 'Total Data')
plt.plot(x,a,color='red', label = 'lambda')
plt.title('lambda = 1/10000')
plt.legend()
plt.grid(True)
plt.show()
```

